

CSC343 Guide

Pierre-William Lessard

Contents

Chapter 1	Relations	Page 2
1.1	The Relational Model	2
1.2	Relational Algebra	2
	Set Operations on Relations — 7	
1.3	Integrity Constraints	8

Chapter 1

Relations

1.1 The Relational Model

Definition 1.1: Relation

A set of tuples that have a certain number of attributes, or columns.
Each tuple in the relation has the same number of attributes, and each attribute has a unique name.

A relation is defined by schemas which indicate the structure of a relation.
Instances of relations include tuples which all follow a relation's schema.

Relations are traditionally thought of as tables. Where the rows are denoted as **tuples** and the columns as “attributes”. The “arity” of a relation is its number of attributes, and its “cardinality” is the number of tuples within a relation instance.

In Relational Algebra, a relation is a set of tuples, meaning:

- There are no duplicate tuples
- The order of tuples does not matter

Definition 1.1.1: Key

A key for a relation, R , is defined as a set of attributes a_1, a_2, \dots, a_n such that:

- \forall tuples $t_1, t_2 \in R$. $\exists i \in [1, n] \cap \mathbb{N}$. $t_1[a_i] \neq t_2[a_i]$
This means that no two tuples may have all the attributes denoted in the key, agree.
- There is no subset of a_1, a_2, \dots, a_n with this same property

This means that keys uniquely identify tuples in a relation's instance, as there must not exist another tuple with the same values for the attributes indicated by the key.

Keys are denoted in the schema by underlining the attributes of the key. A **superkey** is any set of attributes which meet the first property defined above, yet not necessarily the second.

1.2 Relational Algebra

Relational algebra is a type of algebra where the operands are instances of relations. It is a formal system for manipulating and querying data stored in relational databases. It is based on a set of fundamental operations that allow us to combine data from multiple tables in a declarative way. Some of the most common operations

in relational algebra include selection, projection, union, intersection, and difference. These operations allow us to specify the data we want to retrieve from a database, as well as the relationships between different data sets. Relational algebra is an important theoretical foundation for the design and implementation of database systems.

Projection

Projection refers to the operation that extracts certain columns from a relation (table). For example, if you have a relation with columns "name", "age", "gender", and "country", and you only want to see the "name" and "age" columns, you would use the projection operation to create a new relation that only contains those columns. This operation is useful for simplifying and organizing the information in a relation, and for focusing on the specific data that you are interested in.

Definition 1.2.1: Projection

Projection is denoted with the π symbol. Let R be an arbitrary relation with some subset of attributes, a_1, a_2, \dots, a_n . Projecting a_1, a_2, \dots, a_n onto R is notated as:

$$\pi_{a_1, a_2, \dots, a_n} R$$

The result is a new relation with all the same tuples as R , but only including attributes a_1, a_2, \dots, a_n

Selection

Selection is an operation that selects a subset of tuples from a relation. It is typically denoted by the symbol σ , and it is typically used in combination with other operations such as projection and soon to be learned join. Selection is a fundamental operation in relational algebra and is used to specify the desired subset of a relation in order to answer a query. For example, consider a relation R with attributes A , B , and C , and suppose we want to select all tuples from R where the value of attribute A is greater than 5. We can use the select operator to return a relation where this is true.

Definition 1.2.2: Selection

Selection is denoted with the σ symbol. Let R be an arbitrary relation with some subset of attributes, a_1, a_2, \dots, a_n . Let $L(a_1, a_2, \dots, a_n)$ be a logical statement incorporating a subset attributes of R . Selecting tuples which render $L(a_1, a_2, \dots, a_n)$ true in R is notated as:

$$\sigma_{L(a_1, a_2, \dots, a_n)} R$$

The result is a new relation with all the same tuples as R , but only including tuples where $L(a_1, a_2, \dots, a_n)$ rendered true.

Overall, the selection operation is a powerful tool for filtering and selecting subsets of data from a relation, based on specified conditions. It is often used in combination with other operations in relational algebra, such as projection and join, to manipulate and query relational data.

Cartesian Product

The Cartesian product is a binary operation that combines every row of one table with every row of another table. The resulting table, known as the Cartesian product, has a number of rows equal to the product of the number of rows in the two input tables. This operation is also known as a simple join.

Definition 1.2.3: Cartesian Product (Join)

Cartesian Product is denoted with the \times symbol. Let R_1 and R_2 be arbitrary relations with some set of attributes A_1 and A_2 respectively. Joining tuples R_1 and R_2 is notated as:

$$R_1 \times R_2$$

The result is a new relation which we will denote R_{\times} . The following defines the new relation:

$$\forall \text{ tuples } t_1 \in R_1. \forall \text{ tuples } t_2 \in R_2. \exists \text{ a tuple } t_{\times} \in R_{\times}. t_1 | t_2 = t_{\times}$$

Where $t_1 | t_2$ is the tuple such that the attributes of either tuple is concatenated.

This means that each tuple in the first relation is paired with every tuple in the next relation. Therefore, the new table includes all possible combinations of rows from the two input tables. Meaning, given two relations R_1 and R_2 with cardinalities n_1 and n_2 , the cardinality of $R_1 \times R_2$ is $n_1 \cdot n_2$. The arities of the resulting relation is the sum of the two relations' original arities.

Example 1.2.1

Let's say we have two relations, Students and Courses, with the following schemas:

- Students(ID, Name, Age, Gender)
- Courses(Code, Name, Credits)

The cartesian product of these two relations would be a new relation, StudentCourses, with the schema: StudentCourses(ID, Name, Age, Gender, Code, CourseName, Credits)

The relation would have a row for every possible combination of a student and a course, with the student's ID, name, age, and gender paired with the course's code, name, and credits.

Let us say we have an instance of relations Students and Courses, denoted by S and C :

ID	Name	Age	Gender	Code	CourseName	Credits
1	Alice	18	F	CS101	Introduction to Computer Science	4
2	Bob	19	M	MATH201	Calculus I	4
...	ENGL101	Composition	3
...

The result of $S \times C$ would be the following:

ID	Name	Age	Gender	Code	CourseName	Credits
1	Alice	18	F	CS101	Introduction to Computer Science	4
1	Alice	18	F	MATH201	Calculus I	4
1	Alice	18	F	ENGL101	Composition	3
2	Bob	19	M	CS101	Introduction to Computer Science	4
2	Bob	19	M	MATH201	Calculus I	4
2	Bob	19	M	ENGL101	Composition	3
...

This table shows every possible combination of a student and a course, with each row representing a student taking a specific course. Note that the original relation schemas are repeated in the new schema to indicate that each attribute belongs to a specific relation.

Note:-

There exists another join operation, named the “Theta join”. This operation is denoted by $R_1 \bowtie_{L(A)} R_2$ where $L(A)$ is some logical statement asserting something about the attributes of the relations being joined. This join is simply shorthand for $\sigma_{L(A)}(R_1 \times R_2)$.

Natural Join

A natural join is a way of combining two relations into a single relation. The resulting relation includes only pairs of tuples whose values of similar attributes agree. Natural joins are useful when combining two relations which agree on some attribute.

Definition 1.2.4: Natural Join

Natural Join is denoted with the \bowtie symbol. Let R_1 and R_2 be arbitrary relations with some set of attributes A_1 and A_2 respectively. Assume there exists some attribute a that is in both A_1 and A_2 . Natural joining tuples R_1 and R_2 is notated as:

$$R_1 \bowtie R_2$$

The result is a new relation. Let A be the set of attributes that agree between A_1 and A_2 , meaning $A = A_1 \cap A_2$. The following defines the new relation in terms of select (σ) and cartesian product (\times):

$$R_1 \bowtie R_2 = \sigma_{\forall a \in A, R_1[a]=R_2[a]}(R_1 \times R_2)$$

Meaning the natural join of two relations is simply their cartesian product, where the pairs selected agree on the shared attributes. In the final relation, the shared attributes are expressed as one, since we already know they agree.

This property of natural join makes it very useful when combining two relations who are defined as agreeing on some attribute (which we will explain in the next chapter on integrity constraints). The following example illustrates this on two relations, $R(A, B, C)$ and $S(C, D, E)$. Note that these relations agree on the attribute ‘C’.

Example 1.2.2

R			S			R \bowtie S				
A	B	C	C	D	E	A	B	C	D	E
1	2	3	3	7	8	1	2	3	7	8
4	5	6	6	9	10	4	5	6	9	10

This example shows two tables R and S, which are joined using the natural join operation. The resulting table contains all combinations of rows from R and S where the values in the C column match.

Important properties of Natural Join are as follows:

- Natural join is associative, *i.e.* for all arbitrary relations R, S, T : $R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$
- Natural join is commutative, *i.e.* $R \bowtie S = S \bowtie R$
- Natural join can produce unintended results when there are attributes of the same name which are not meant to be related.
- When no attributes match between the two relations, natural join acts the same as cartesian product

Assignment and Rename

When operating with many relations and results of operations it is helpful to assign new names to said results. By assigning names to results we are able to recall previous operations with limited redundancy. The assignment operation is used in this case. The operation uses the $:=$ operator.

Definition 1.2.5: Assignment

For any given relational expression, E , with resulting attributes n resulting attributes, an assignment is as follows:

$$R(A_1, \dots, A_n) := E$$

Where R is the name of the new assigned relation, and A_1, \dots, A_n are the new (or original) names of the attributes.

In multi-step operations of relational algebra, the assignment operation saves space and time rewriting redundant operations. Renaming attributes is important when joining two instances of the same relation. An important thing to note is that the name of the new assigned relation R , must not be one of the relations in the schema. Assignment is **not** equivalent or used to update relations.

It is also possible to assign names to an instance of a relation in an operation. This is different from what the assignment operations does, which is save a result for later use. The rename operator, using the greek symbol ρ , rho, is meant to alleviate ambiguity when relational algebra references an attribute that may come from multiple results.

Example 1.2.3

Let R be a relation with schema $R(A, B)$.

For this example R will have two tuples $(1, 1), (0, 0)$. Observe the following operation:

$$\sigma_{A=B}(R \times R)$$

If A references the attribute from the left hand side of the product, and B references the right, then the resulting relation is empty. However, if both variables reference the left side then the resulting relation has the maximum cardinality of 4.

This example is simple yet demonstrates the potential ambiguity that may be present during a more complicated combination of operations. For this reason we have the renaming operator

Definition 1.2.6: Rename

To rename an instance of a relation, inside of an expression, the ρ symbol can be used with a new name referenced in the subscript. Any following usage of an attribute preceded by the new name and a dot will reference the attributes of the renamed instance. Note this is only for the expression where the relation is renamed. Renaming any relation R , with name T , is denoted as

$$\rho_T R$$

Specifying attributes, is possible through following T with new attribute names enclosed in parenthesis.

Example 1.2.4

Following the previous example, the ambiguity can be removed by specifying a name for each relation.

$$\sigma_{R_1.A=R_2.B}(\rho_{R_1} R \times \rho_{R_2} R) \tag{1.1}$$

$$\sigma_{R_1.A=R_1.B}(\rho_{R_1} R \times \rho_{R_2} R) \tag{1.2}$$

The operation in (1.1) now definitely produces the empty relation, $\{\}$, and the operation in (1.2) produced the every possible combination from $R \times R$

1.2.1 Set Operations on Relations

The final important operation on relations directly follows from relations being defined as set adjacent. This means that the main set operations are possible and very applicable to relations.

Set operations such as union, intersection, and difference can be used in relational algebra. Set operations allow you to combine two or more relations to form a new relation, based on the criteria specified in the set operation. A caveat is that the two relations being operated on must agree on the number, name, and order of attributes. This ensures the rules of set relations directly apply.

Definition 1.2.7: Relational Algebra Set Operations

Given two relations R_1 and R_2 , where the attributes of R_1 and R_2 agree in order, name, and number (the schemas are the same), then the following set operations from set theory are applicable:

- Union: \cup
- Intersect: \cap
- Difference: $-$

Example 1.2.5

Suppose we have two relations, R and S, with the following attributes and tuples:

A	B	C	A	B	C
1	2	3	1	4	7
4	5	6	2	5	8
7	8	9	3	6	9

The operation, $R \cup S$, results in the following table:

A	B	C
1	2	3
4	5	6
7	8	9
1	4	7
2	5	8
3	6	9

As is obvious from the example, the relations must have equal attributes. Otherwise, a set operation is not applicable. The same can be seen for the similar example of the intersect set operation:

Example 1.2.6

A	B		A	B					
1	2	∩	3	4	=				
3	4		5	7					
5	6		9	8					
			<table><tr><td>A</td><td>B</td></tr><tr><td>3</td><td>4</td></tr></table>			A	B	3	4
A	B								
3	4								

In this example, we have two relations, A and B, and the resulting table after performing the intersect operation contains only the tuple (3, 4), since it is the only tuple that appears in both of the original tables.

1.3 Integrity Constraints

Referential integrity constraints are important in constraining more custom rules onto a relation. They sometimes refer to some reference between relations. This can be simple as an instance's attribute's values needing to be a subset of some other relations. This can become much more complicated through using relational algebra.

Referential Integrity Constraints

A referential integrity constraint is a rule that is used to ensure that relationships between tables in a database are maintained. It is a way of ensuring that data in one table references data in another table and that the data in the other table is valid and consistent. For example, if a table contains a foreign key that references a primary key in another table, a referential integrity constraint can be used to ensure that the value in the foreign key column is always a valid primary key value in the other table. This helps to maintain the integrity of the data and ensures that the relationships between the tables are consistent. More rigorously we can define it as follows:

Definition 1.3.1: Referential Integrity Constraint

Let R_1 and R_2 be two unique relations. Let X and Y be arbitrary sets of attributes of equal length in R_1 and R_2 respectively. The referential integrity constraint between these two relations, denoted $R_1[X] \subseteq R_2[Y]$ implies:

$$\forall \text{ values } v_1 \in R_1[X]. \exists \text{ a value } v_2 \in R_2[Y]. v_1 = v_2$$

Meaning the set of values in $R_1[X]$ are a subset of the values in $R_2[Y]$

Referential Integrity Constraints are important in defining the relationship between two relations. When Y is a key we call the referential integrity constraint a **Foreign Key Constraint**.